

Title: Robust Regression Engine

Duration: 6 Hours

Objective

The objective of this project is to evaluate students' understanding of advanced supervised learning regression techniques, with a strong focus on regularization, model generalization, cross-validation strategies, and tree-based regression algorithms. Students will learn how to control overfitting, select optimal models, and compare linear vs non-linear regressors using real-world data.

Problem Statement

You are working as a **Machine Learning Engineer** at a real estate analytics company. The company already uses a basic regression model to predict **house prices**, but it suffers from **overfitting and unstable predictions** across different datasets.

Your task is to build a **robust regression pipeline** that:

- Applies **regularization techniques**
- Uses **proper cross-validation strategies**
- Compares **linear and non-linear regression models**
- Selects the **best-performing model** based on validation performance

The final solution must generalize well on unseen data.

Part A: Conceptual Foundation (Theory)

Answer briefly (2–4 lines each):

1. What is **Regularization** in Machine Learning? Why is it needed?
 2. Difference between **Ridge Regression (L2)** and **Lasso Regression (L1)**.
 3. What is **Cross-Validation** and why is it important?
 4. Explain the following cross-validation techniques:
 - o K-Fold Cross-Validation
 - o Stratified K-Fold Cross-Validation
 - o Leave-One-Out Cross-Validation (LOOCV)
 - o Time Series Split
 5. Why are **tree-based models** less sensitive to feature scaling?
-



Part B: Dataset Understanding & Preparation

You are provided with a **real estate dataset** containing:

- Property size and structural attributes
- Location-based numerical indicators
- Property age and renovation details
- Target variable: **House Price**

Dataset: [Access from here](#)

Tasks:

6. Identify **features and target variable**.
 7. Perform a **train-test split**.
 8. Apply **basic preprocessing** (scaling where required).
-



Part C: Regularized Linear Models

9. Implement **Ridge Regression (L2)**.
10. Implement **Lasso Regression (L1)**.
11. Tune the **regularization parameter (α)** using **cross-validation**.
12. Compare Ridge and Lasso based on:
 - Training error
 - Validation error
 - Feature coefficient behavior

Part D: Cross-Validation Strategies

13. Apply and compare the following cross-validation techniques:
 - K-Fold Cross-Validation
 - Stratified K-Fold Cross-Validation (*by binning target values*)
 - Leave-One-Out Cross-Validation
 - Time Series Split (*using a time-related feature*)
 14. Analyze how performance metrics vary across different CV strategies.
-

Part E: Tree-Based Regression Models

15. Implement **Decision Tree Regression**.
 16. Control tree complexity using hyperparameters (max depth, min samples).
 17. Implement **Random Forest Regression**.
 18. Compare single-tree vs ensemble performance.
-

Part F: Support Vector Regression

19. Implement **Support Vector Regression (SVR)** with:
 - Linear kernel
 - Polynomial or RBF kernel
 20. Tune hyperparameters (C , γ , ϵ).
 21. Compare SVR performance with linear and tree-based models.
-

Part G: Model Comparison & Evaluation

22. Evaluate all models using appropriate regression metrics:
 - MSE, MAE, RMSE
 - R^2 Score
23. Compare:
 - Regularized Linear Models
 - Tree-Based Models

- Support Vector Regression
24. Identify signs of **overfitting or underfitting** in each model.
-

Part H: Final Analysis & Reporting

25. Prepare a final report covering:
- Best-performing model and justification
 - Impact of regularization
 - Role of cross-validation in model stability
 - Comparison of linear vs non-linear regressors
 - Business interpretation of results
26. Submit:
- Source code / notebook
 - Evaluation tables and plots
 - Final conclusions
-

Submission Guidelines

- Include practical implementation in a Jupyter Notebook and screenshots.
- Label all charts clearly and write short interpretations under each result.
- **GitHub Repository:**
 - Create a GitHub repository to host your project.
 - Upload your project files, including source code, and documentation to the repository.
 - Add a document (PDF) explaining theory concepts with definitions (if any).
 - Ensure that you provide a clear and descriptive README.md file.

Remember to follow the instructions provided professionally, make suitable assumptions wherever necessary, and avoid copying code or content from unauthorized sources. Good luck with your project work!

Robust Regression Engine
Supervised Learning

BRING ON YOUR CODING ATTITUDE